

Lecture 1:

Why Parallelism? Why Efficiency?

Parallel Computer Architecture and Programming
CMU 15-418/15-618, Spring 2020

Hi!



Randy Bryant

Plus . . .

An evolving collection of teaching assistants



Nathan Beckmann

Getting into the Class

- **Status (Mon Jan. 13, 09:30)**
 - 157 students enrolled
 - 103 on wait list
 - 175 max. enrollment
 - ~20 slots
- **If you are registered**
 - **Do Assignment 1**
 - Due Jan. 29
 - **If find too challenging, then please drop by Jan. 27**
- **Clearing Wait List**
 - **Complete Assignment 1 by Jan. 22, 23:00**
 - **No Autolab account required**
 - **We will enroll top-performing students**
 - ***It's that simple!***
 - **You will know by Jan. 27**

What will you be doing in this course?

Assignments

- **Four programming assignments**

- **First assignment is done individually, the rest will be done in pairs**
 - **Each uses a different parallel programming environment**
 - **Each also involves measurement, analysis, and tuning**



Assignment 1: SIMD and multi-core parallelism



Assignment 2: CUDA programming on NVIDIA GPUs



Assignment 3: Parallel Programming via a Shared-Address Space Model



Assignment 4: Parallel Programming via a Message Passing Model

Final project

- **6-week self-selected final project**
- **Performed in groups (by default, 2 people per group)**
- **Keep thinking about your project ideas starting TODAY!**
- **Poster session at end of term**
- **Check out previous projects:**

<http://15418.courses.cs.cmu.edu/spring2016/competition>

<http://15418.courses.cs.cmu.edu/fall2017/article/10>

<http://www.cs.cmu.edu/afs/cs.cmu.edu/academic/class/15418-s18/www/15418-s18-projects.pdf>

<http://www.cs.cmu.edu/afs/cs.cmu.edu/academic/class/15418-s19/www/15418-s19-projects.pdf>

Exercises

- **Five homework exercises**
 - **Scheduled throughout term**
 - **Designed to prepare you for the exams**
 - **We will grade your work only in terms of participation**
 - **Did you make a serious attempt?**
 - **Only a participation grade will go into the gradebook**

Grades

40% Programming assignments (4)

30% Exams (2)

25% Final project

5% Exercises

Each student gets up to five late days on programming assignments (see syllabus for details)

Getting started

- Visit course home page
 - <http://www.cs.cmu.edu/~418/>
- Sign up for the course on Piazza
 - <http://piazza.com/cmu/spring2020/1541815618>
- Textbook
 - There is no course textbook, but please see web site for suggested references
- Find a Partner
 - Assignments 2–4, final project

Regarding the class meeting times

- **Class MWF 3:00–4:20**
 - **Lectures (mostly)**
 - **Some designated “Recitations”**
 - **Targeted toward things you need to know for an upcoming assignment**
- **No classes last part of the term**
 - **Let you focus on projects**

Collaboration (Acceptable & Unacceptable)

- **Do**
 - **Become familiar with course policy**
<http://www.cs.cmu.edu/~418/academicintegrity.html>
 - **Talk with instructors, TAs, partner**
 - **Brainstorm with others**
 - **Use general information on WWW**
- **Don't**
 - **Copy or provide code to anyone**
 - **Use information specific to 15-418/618 on WWW**
 - **Leave your code in accessible place**
 - **Now or in the future**

A Brief History of Parallel Computing

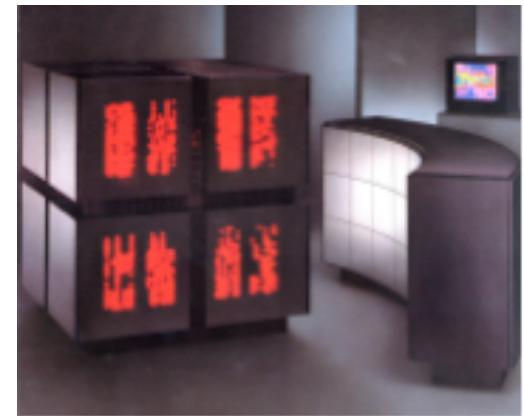
- Initial Focus (starting in 1970s): “Supercomputers” for Scientific Computing



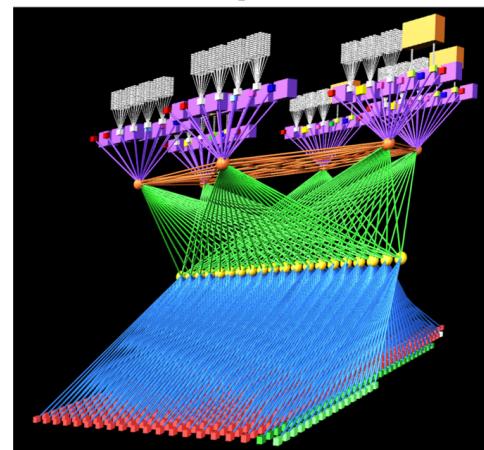
C.mmp at CMU (1971)
16 PDP-11 processors



Cray XMP (circa 1984)
4 vector processors



Thinking Machines CM-2 (circa 1987)
65,536 1-bit processors +
2048 floating-point co-processors



800+ compute nodes
Heterogenous Structure



Bridges at the Pittsburgh
Supercomputer Center

A Brief History of Parallel Computing

- Initial Focus (starting in 1970s): “Supercomputers” for Scientific Computing
- Another Driving Application (starting in early ‘90s): Databases
 - Especially, handling millions of transactions per second for web services



Sun Enterprise 10000 (circa 1997)
16 UltraSPARC-II processors



Oracle Supercluster M7 (today)
4 X 32-core SPARC M2 processors

- RIP 2019. Killed by cloud computing

A Brief History of Parallel Computing

- **Cloud computing (2000–present)**
- **Build out massive centers with many, simple processors**
 - **Connected via LAN technology**
- **Program using distributed-system models**



- **Not really the subject of this course (take 15-440)**

Setting Some Context

- Before we continue our multiprocessor story, let's pause to consider:
 - Q: what had been happening with single-processor performance?
- A: since forever, they had been getting exponentially faster
 - Why?

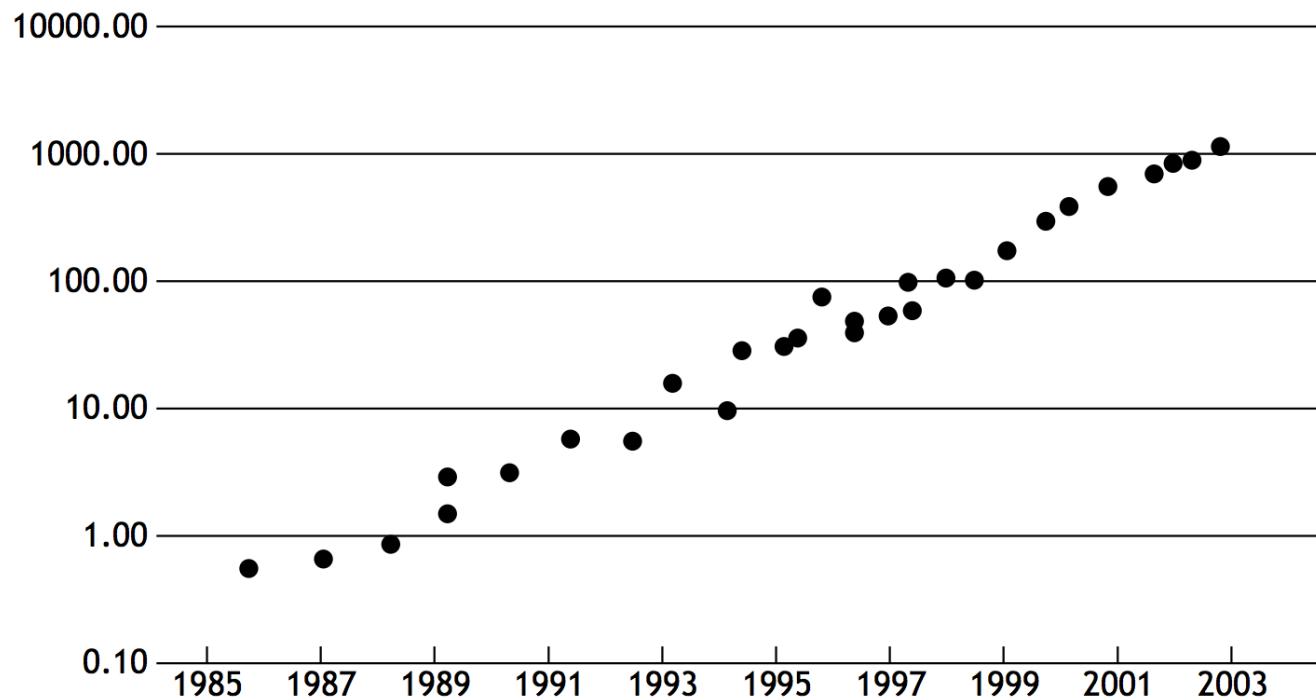


Image credit: Olukutun and Hammond, ACM Queue 2005

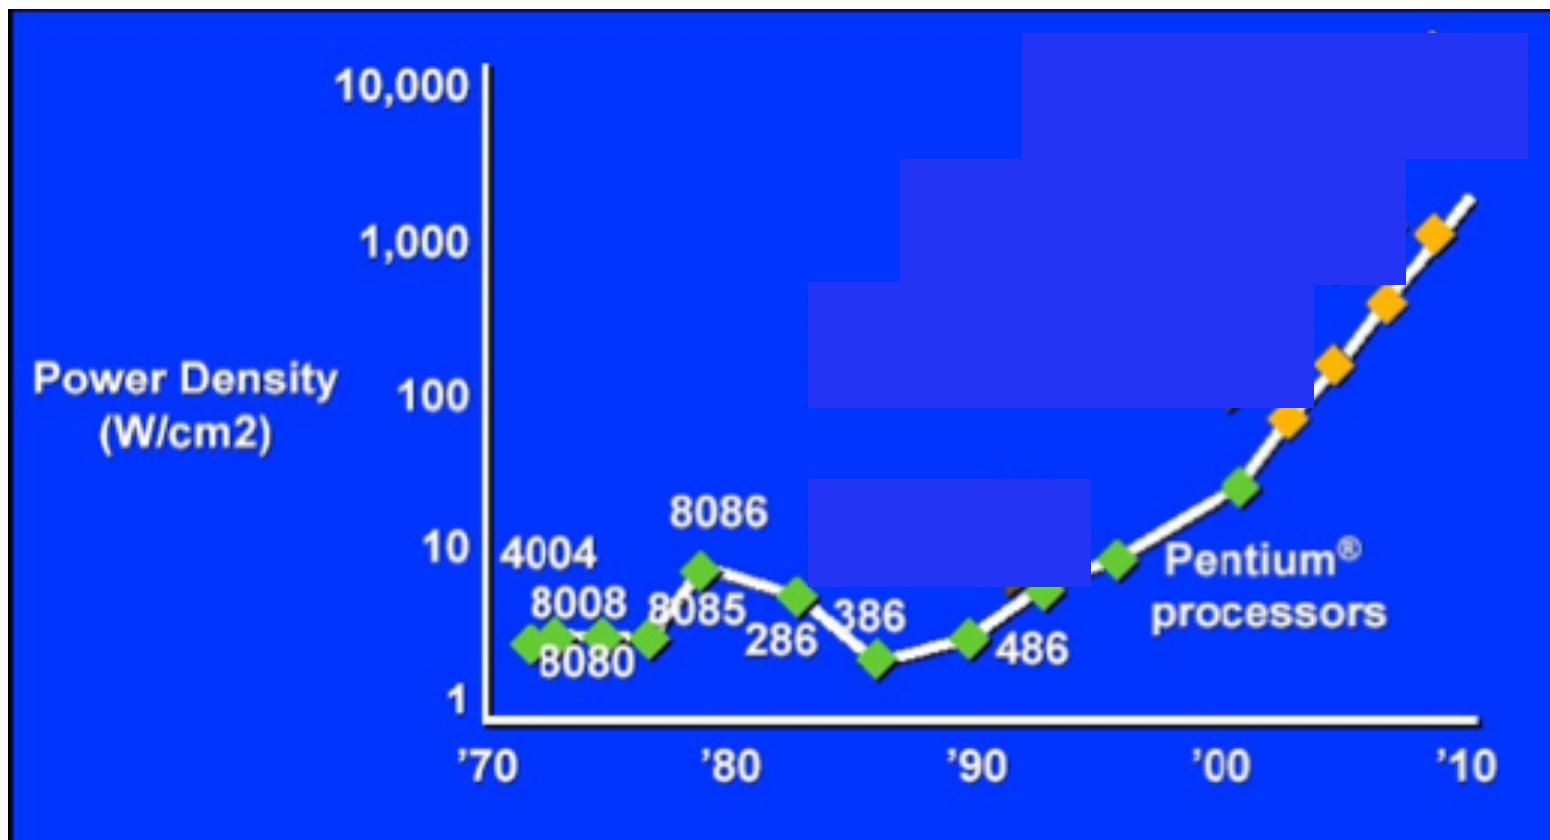
CMU 15-418/618, Spring 2020

A Brief History of Processor Performance

- **Wider data paths**
 - 4 bit → 8 bit → 16 bit → 32 bit → 64 bit
- **More efficient pipelining**
 - e.g., 3.5 Cycles Per Instruction (CPI) → 1.1 CPI
- **Exploiting instruction-level parallelism (ILP)**
 - “Superscalar” processing: e.g., issue up to 4 instructions/cycle
 - “Out-of-order” processing: extract parallelism from instruction stream
- **Faster clock rates**
 - e.g., 10 MHz → 200 MHz → 3 GHz
- **During the 80s and 90s: large exponential performance gains**
 - and then...

A Brief History of Parallel Computing

- Initial Focus (starting in 1970s): “Supercomputers” for Scientific Computing
- Another Driving Application (starting in early ‘90s): Databases
- Inflection point in 2004: Intel hits the Power Density Wall



Pat Gelsinger, ISSCC 2001

CMU 15-418/618, Spring 2020

From the New York Times

Intel's Big Shift After Hitting Technical Wall

The warning came first from a group of hobbyists that tests the speeds of computer chips. This year, the group discovered that the Intel Corporation's newest microprocessor was running slower and hotter than its predecessor.

What they had stumbled upon was a major threat to Intel's longstanding approach to dominating the semiconductor industry - relentlessly raising the clock speed of its chips.

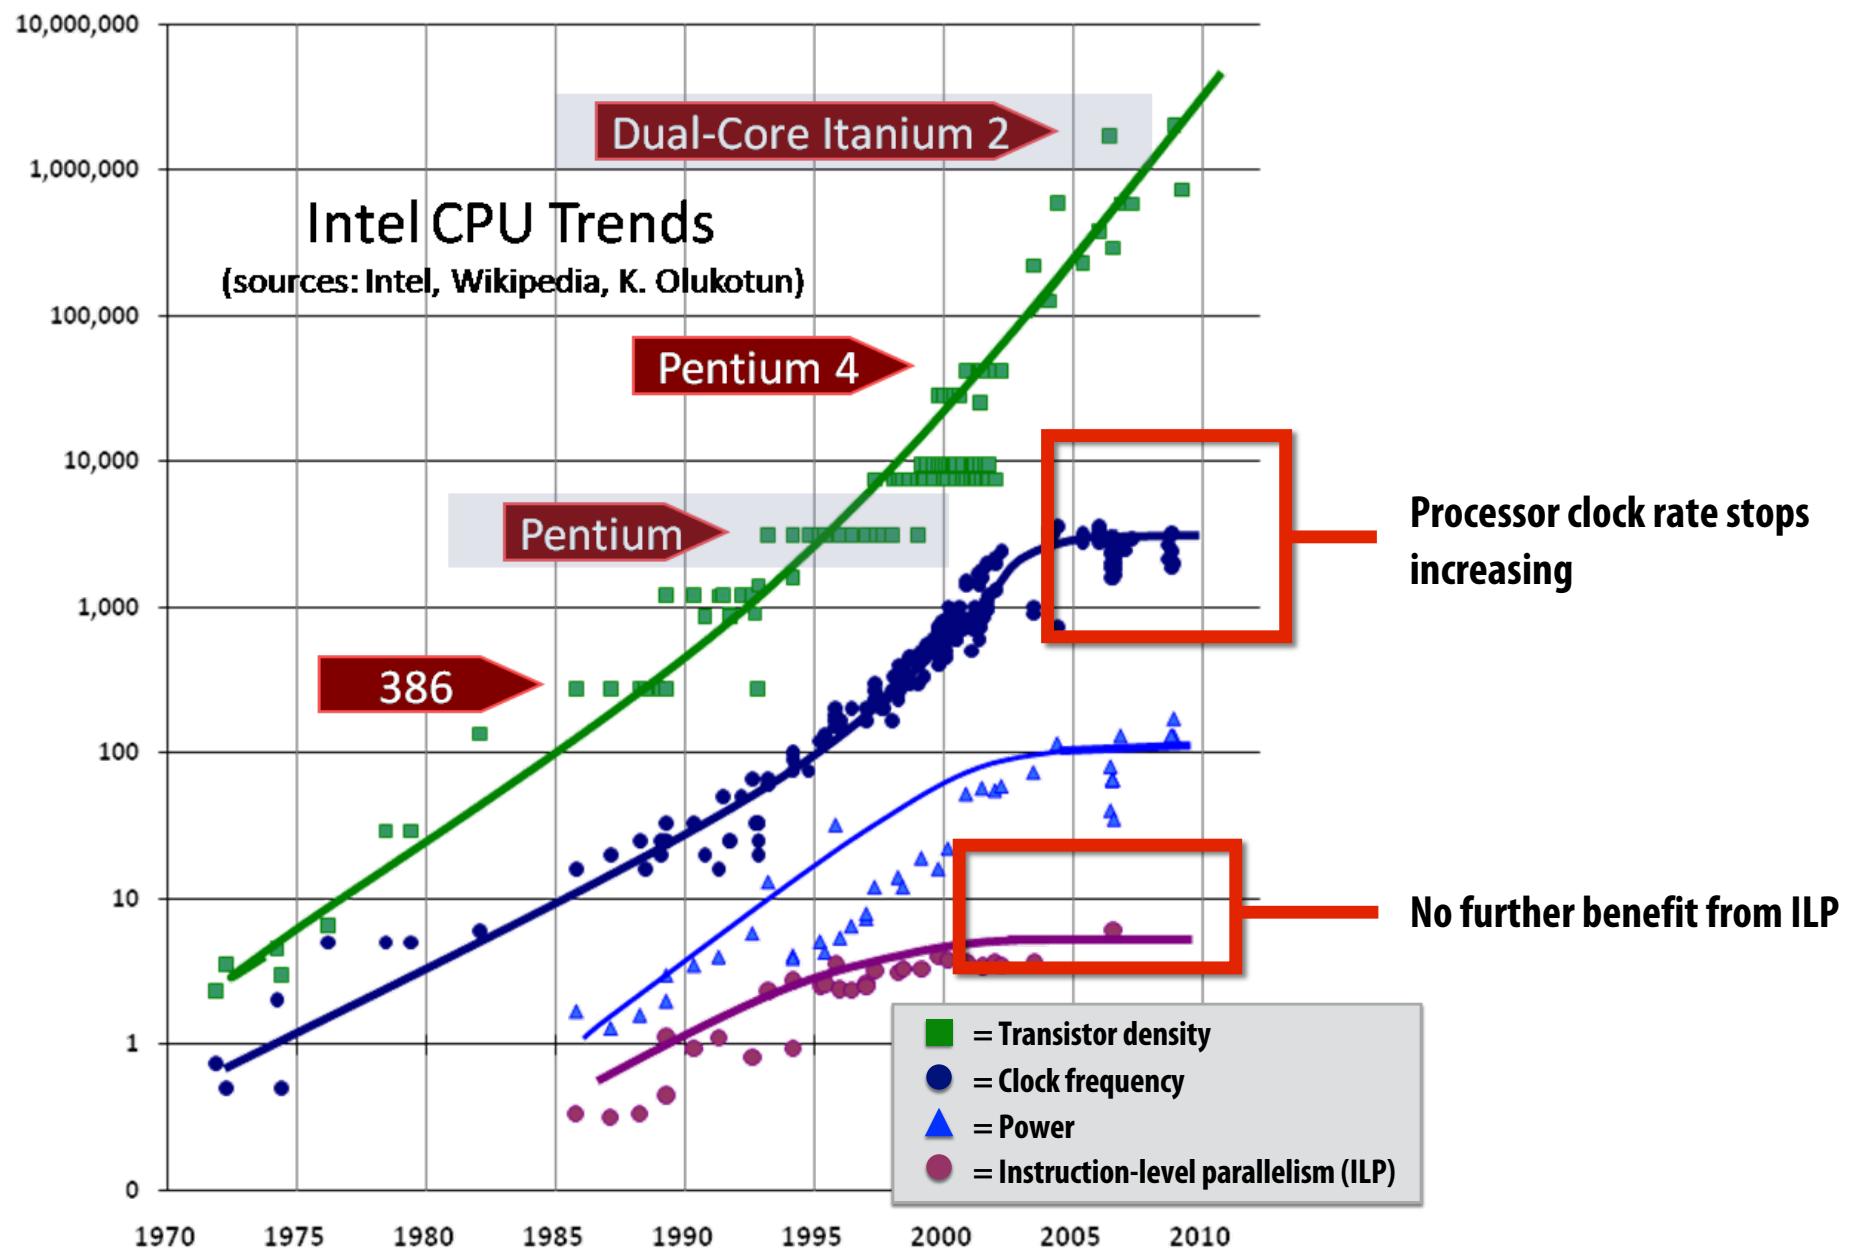
Then two weeks ago, Intel, the world's largest chip maker, publicly acknowledged that it had hit a "thermal wall" on its microprocessor line. As a result, the company is changing its product strategy and disbanding one of its most advanced design groups. Intel also said that it would abandon two advanced chip development projects, code-named Tejas and Jayhawk.

Now, Intel is embarked on a course already adopted by some of its major rivals: obtaining more computing power by stamping multiple processors on a single chip rather than straining to increase the speed of a single processor.

...

John Markoff, New York Times, May 17, 2004

ILP tapped out + end of frequency scaling



Programmer's Perspective on Performance

Question: How do you make your program run faster?

Answer before 2004:

- Just wait 6 months, and buy a new machine!
- (Or if you're really obsessed, you can learn about parallelism.)

Answer after 2004:

- You need to write parallel software.

Parallel Machines Today

Examples from Apple's product line:



Mac Pro
28 Intel Xeon W cores



iMac Pro
18 Intel Xeon W cores



MacBook Pro Retina 15"
8 Intel Core i9 cores

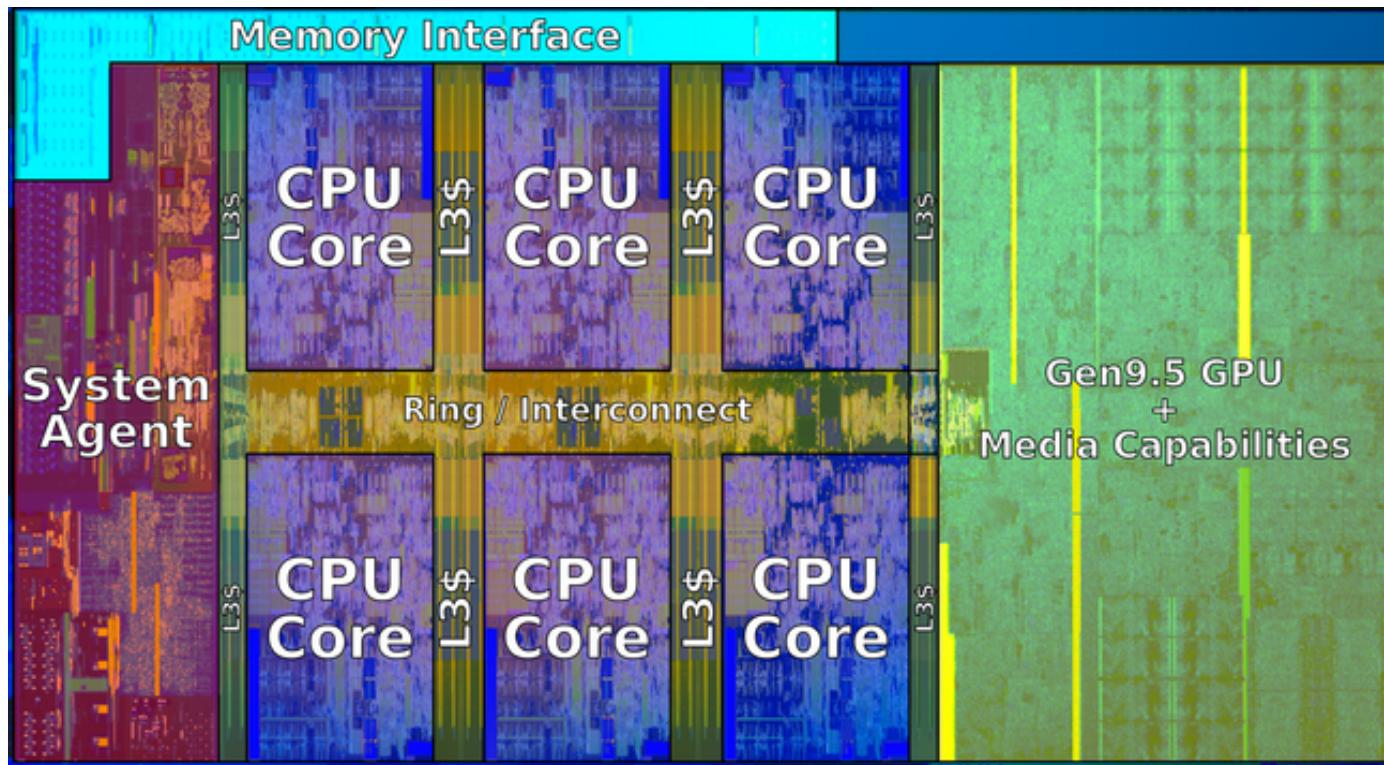


iPhone XS
6 CPU cores
(2 fast +
4 low power)
6 GPU cores

(images from apple.com)

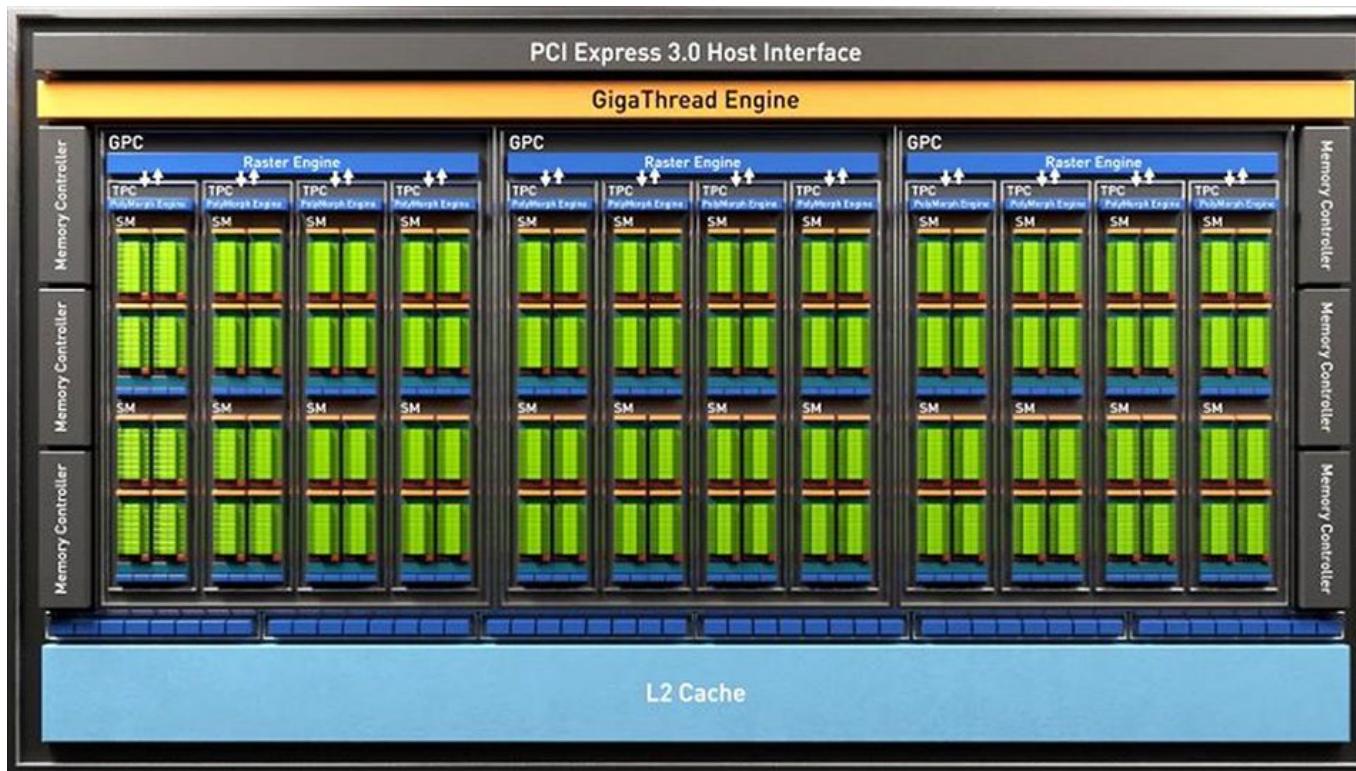
Intel Coffee Lake Core i9 (2019)

6-core CPU + multi-core GPU integrated on one chip



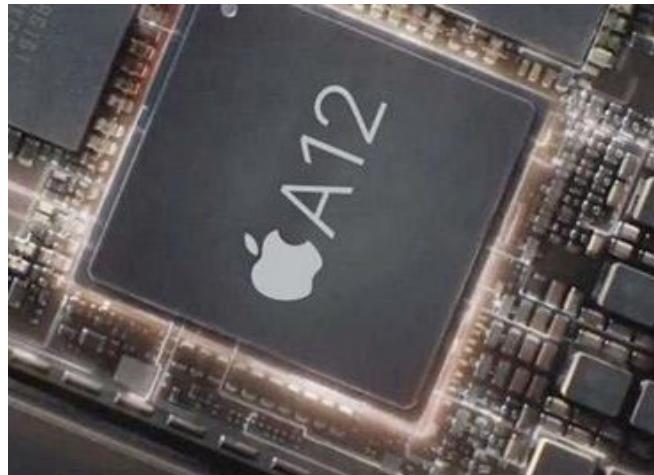
NVIDIA GeForce GTX 1660 Ti GPU (2019)

**24 major processing blocks
(but much, much more parallelism available... details coming soon)**

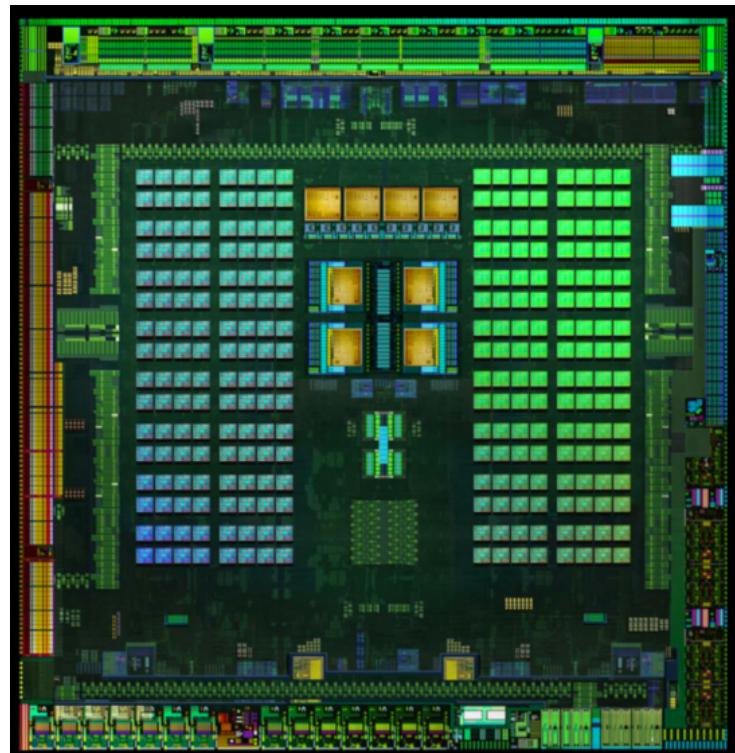


Mobile parallel processing

Power constraints heavily influence design of mobile systems



Apple A12: (in iPhone XR)
4 CPU cores
4 GPU cores
Neural net engine
+ much more



NVIDIA Tegra K1:
Quad-core ARM A57 CPU + 4 ARM A53 CPUs +
NVIDIA GPU + image processor...

Supercomputing

- Today: clusters of multi-core CPUs + GPUs
- Oak Ridge National Laboratory: Summit (#1 supercomputer in world)
 - 4,608 nodes
 - Each with two 22-core CPUs + 6 GPUs



Supercomputers vs. Cloud Systems

Supercomputers

- Few, big tasks

- Customized
- Optimized for reliability
- Low latency interconnect

- Minimal
- Static scheduling

Data Center Clusters

Target Applications

- Many small tasks

Hardware

- Consumer grade
- Optimized for low cost
- Throughput-optimized interconnect

Run-Time System

- Provides reliability
- Dynamic allocation

Application Programming

- Low-level, processor-centric model
- Programmer manages resources

- High level, data-centric model
- Let run-time system manage resources

Supercomputer / Data Center Overlap

- **Supercomputer features in data centers**
 - **Data center computers sometimes used to solve problem**
 - E.g., learn neural network for language translation
 - **Data center computers sometimes equipped with GPUs**
- **Data center features in supercomputers**
 - **Also used to process many small–medium jobs**

What is a parallel computer?

One common definition

A parallel computer is a **collection of processing elements** that **cooperate to solve problems quickly**

We care about performance *
We care about efficiency

We're going to use multiple
processors to get it

* Note: different motivation from “concurrent programming” using pthreads in 15-213

DEMO 1

(This semester's first parallel program)

Speedup

One major motivation of using parallel processing: achieve a speedup

For a given problem:

$$\text{speedup(using } P \text{ processors)} = \frac{\text{execution time (using 1 processor)}}{\text{execution time (using } P \text{ processors)}}$$

Class observations from demo 1

- **Communication limited the maximum speedup achieved**
 - In the demo, the communication was telling each other the partial sums
- **Minimizing the cost of communication improves speedup**
 - Moving students (“processors”) closer together (or let them shout)

DEMO 2

(scaling up to four “processors”)

Class observations from demo 2

- **Imbalance in work assignment limited speedup**
 - Some students (“processors”) ran out work to do (went idle), while others were still working on their assigned task
- **Improving the distribution of work improved speedup**

DEMO 3

(massively parallel execution)

Class observations from demo 3

- The problem I just gave you has a significant amount of communication compared to computation
- Communication costs can dominate a parallel computation, severely limiting speedup

Course theme 1:

Designing and writing parallel programs ... that scale!

- **Parallel thinking**
 1. Decomposing work into pieces that can safely be performed in parallel
 2. Assigning work to processors
 3. Managing communication/synchronization between the processors so that it does not limit speedup
- **Abstractions/mechanisms for performing the above tasks**
 - Writing code in popular parallel programming languages

Course theme 2:

Parallel computer hardware implementation: how parallel computers work

- **Mechanisms used to implement abstractions efficiently**
 - Performance characteristics of implementations
 - Design trade-offs: performance vs. convenience vs. cost
- **Why do I need to know about hardware?**
 - Because the characteristics of the machine really matter
(recall speed of communication issues in earlier demos)
 - Because you care about efficiency and performance
(you are writing parallel programs after all!)

Course theme 3:

Thinking about efficiency

- **FAST != EFFICIENT**
- Just because your program runs faster on a parallel computer, it does not mean it is using the hardware efficiently
 - Is 2x speedup on computer with 10 processors a good result?
- Programmer's perspective: make use of provided machine capabilities
- HW designer's perspective: choosing the right capabilities to put in system (performance/cost, cost = silicon area?, power?, etc.)

Fundamental Shift in CPU Design Philosophy

Before 2004:

- within the chip area budget, maximize **performance**
 - **increasingly aggressive speculative execution for ILP**

After 2004:

- **area within the chip matters** (limits # of cores/chip):
 - **maximize performance per area**
- **power consumption** is critical (battery life, data centers)
 - **maximize performance per Watt**
- **upshot:** major focus on **efficiency** of cores

Summary

- **Today, single-thread performance is improving very slowly**
 - **To run programs significantly faster, programs must utilize multiple processing elements**
 - **Which means you need to know how to write parallel code**
- **Writing parallel programs can be challenging**
 - **Requires problem partitioning, communication, synchronization**
 - **Knowledge of machine characteristics is important**
- **I suspect you will find that modern computers have tremendously more processing power than you might realize, if you just use it!**
- **Welcome to 15-418!**