

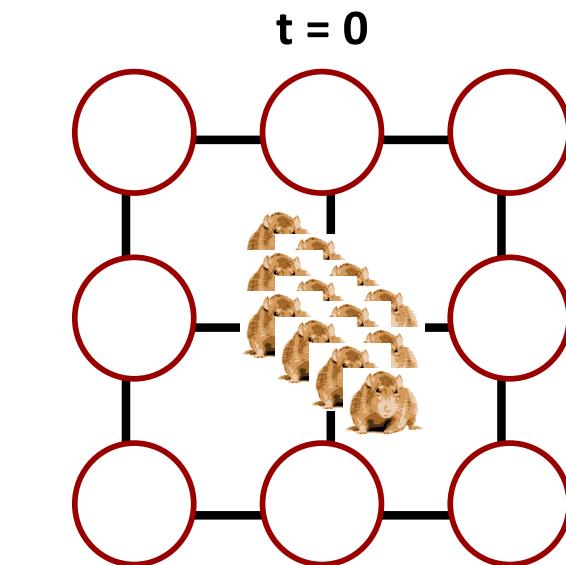
# Assignment 3: GraphRats



# Topics

- Application
- Implementation Issues
- Optimizing for Parallel Performance
- Useful Advice

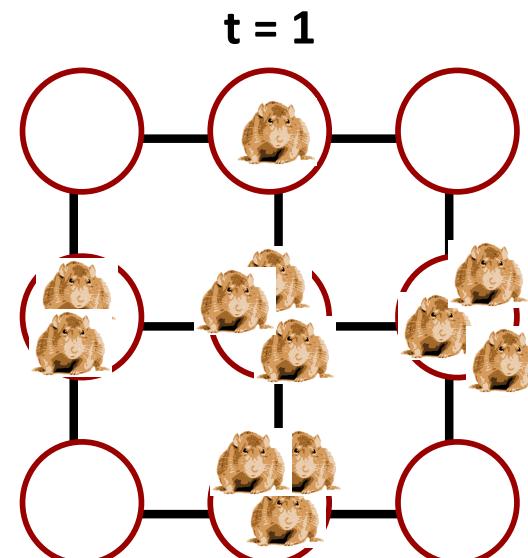
# Basic Idea



- **Graph**
  - $W \times H$  grid
- **Initial State**
  - Start with all  $R$  rats in center

## ■ Transitions

- Each rat decides where to move next
  - Don't like crowds
  - But also don't like to be alone
- Weighted random choice



# Node Count Representation (18 x 12)

# Simulation Example (18 x 12)

- Note moves to nodes not connected by grid
    - Explanation to follow

# Simulation Example (18 x 12)

t = 20.

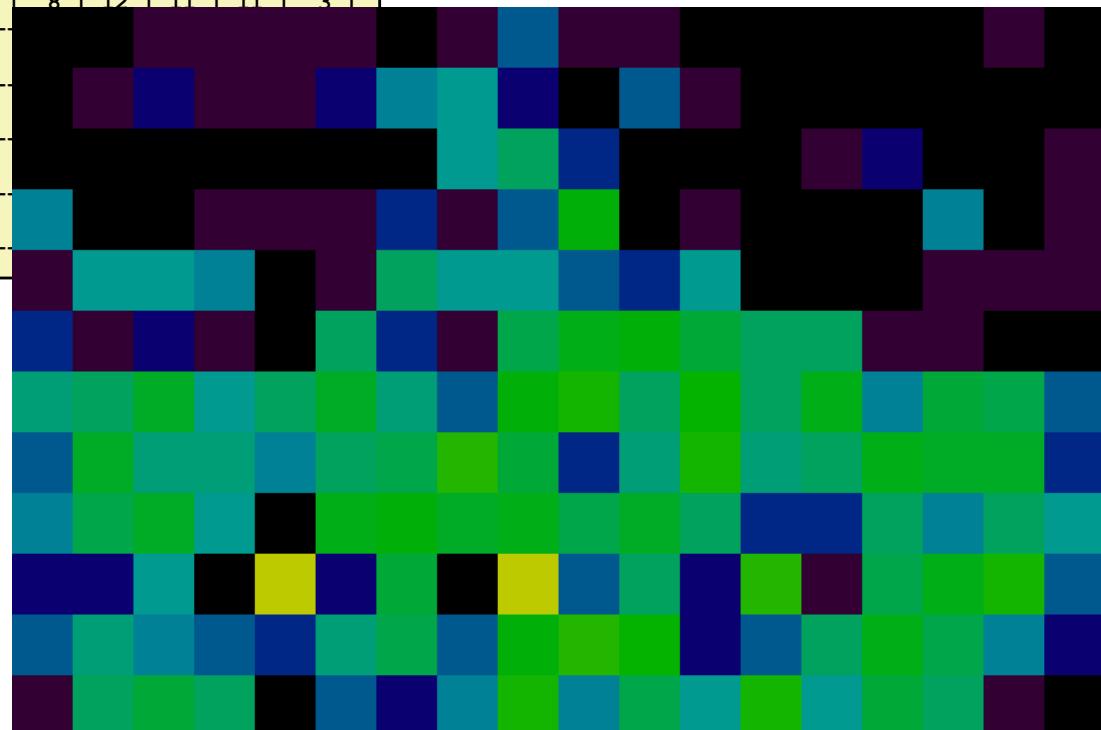
			1		1		1			1		4		1		1						1														
			1		2		1		1		2		5		6		2			4		1														
													6		8		3					1		2		1										
	5				1		1		1		3		1		4		13			1				5		1										
	1		6		6		5			1		8		6		6		4		3		6				1		1								
	3		1		2		1			8		3		1		9		12		13		10		8		8		1		1						
	7		8		11		6		8		11		7		4		13		15		8		14		8		12		5		10		9		4	
	4		11		7		7		5		8		9		16		10		3		7		15		7		8		12		11		11		3	
	5		9		11		6			12		13		11		12		9		11		8		3		3		8		5		8		6		
	2		2		6			31		2		10			31		4		8		2		16		1		9		12		15		4			
	4		7		5		4		3		7		9		4		13		16		14		2		4		8		12		9		5		2	
	1		8		10		8			4		2		5		15		5		9		6		15		6		10		8		1				

## ■ Rats dispersed across graph

# Visualizations

t = 20.																			
			1	1	1	1		1	4	1	1	1					1		
		1	2	1	1	2	5	6	2		4	1							
							6	8	3				1	2			1		
	5			1	1	1	3	1	4	13		1				5		1	
	1	6	6	5		1	8	6	6	4	3	6				1	1	1	1
	3	1	2	1		8	3	1	9	12	13	10	8	8	1	1			
	7	8	11	6	8	11	7	4	13	15	8	14	8	12	5	10	9	4	
	4	11	7	7	5	8	9	16	10	3	7	15	7	8	12	11	11	3	
	5	9	11	6		12	13	11	12	9	11	8	3						
	2	2	6		31	2	10		31	4	8	2	16						
	4	7	5	4	3	7	9	4	13	16	14	2	4						
	1	8	10	8		4	2	5	15	5	9	6	15						

## Text (“a” for ASCII)



## Heat Map (“h”)

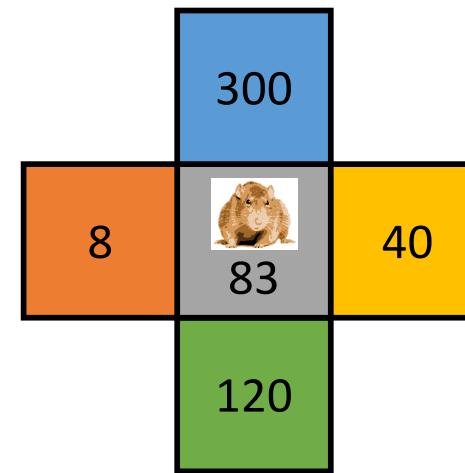
# Running it yourself

```
linux> cd some directory
linux> git clone https://github.com/cmu15418/asst3-s20.git
linux> cd asst3-s20/code
Linux> make demoX
    X from 1 to 11
```

## ■ Demos

- 1: Text visualization, synchronous updates
- 2: Heap-map, synchronous updates

# Determining Rat Moves

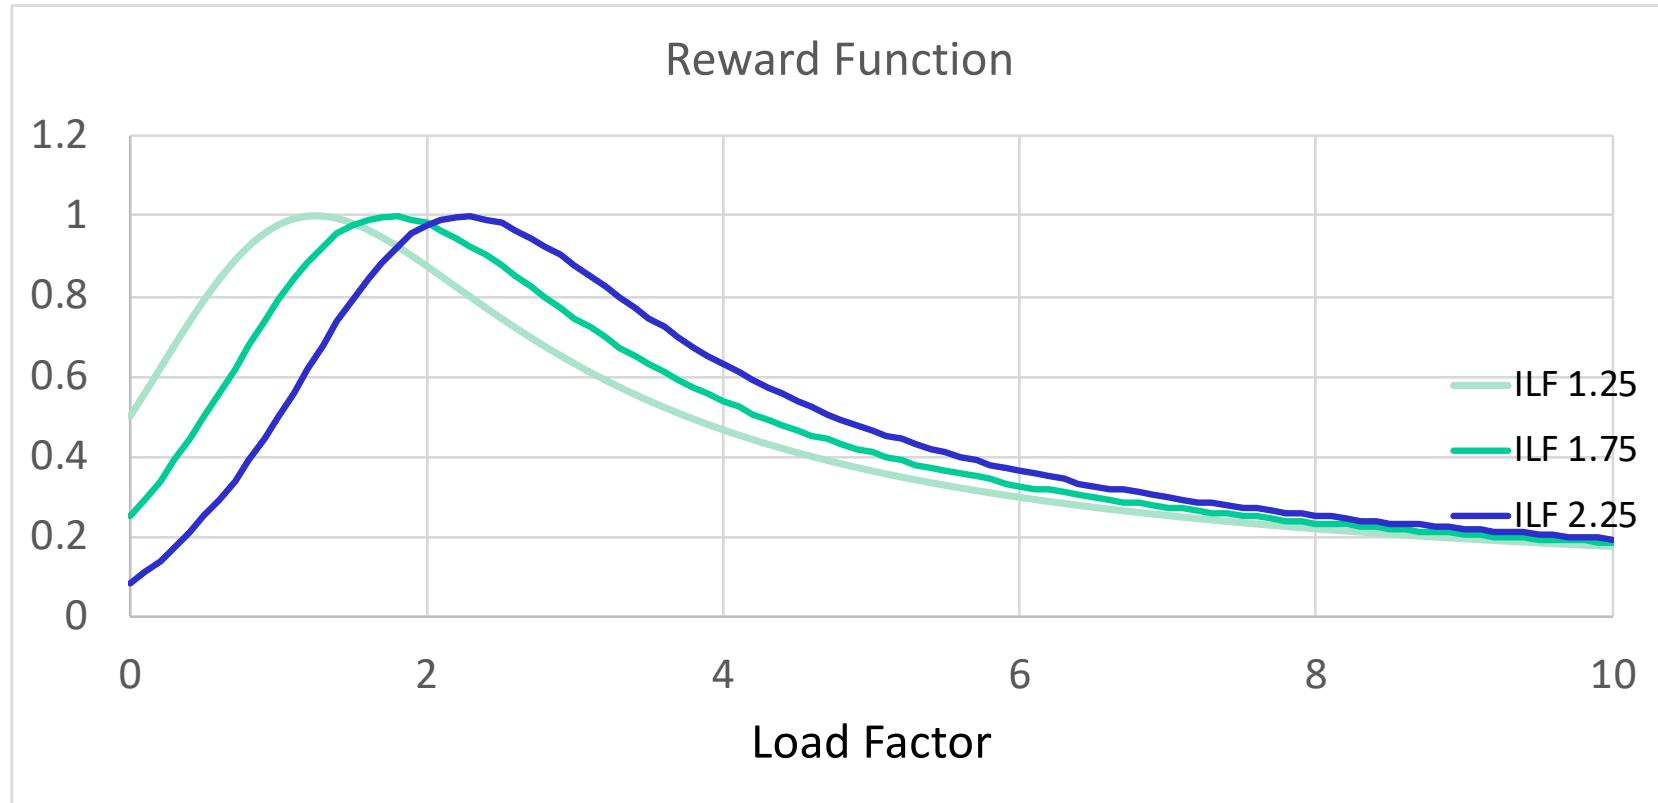


- **Count number of rats at current and adjacent locations**
  - Adjacency structure represented as graph
- **Compute reward value for each location**
  - Based on *load factor*  $l$  = count/average count
  - $l^*$       Ideal load factor (ILF) (varying)
  - $\alpha$       Fitting parameter (= 0.4)

$$\text{Reward}(l) = \frac{1}{1 + (\log_2 [1 + \alpha(l - l^*)])^2}$$

# Reward Function

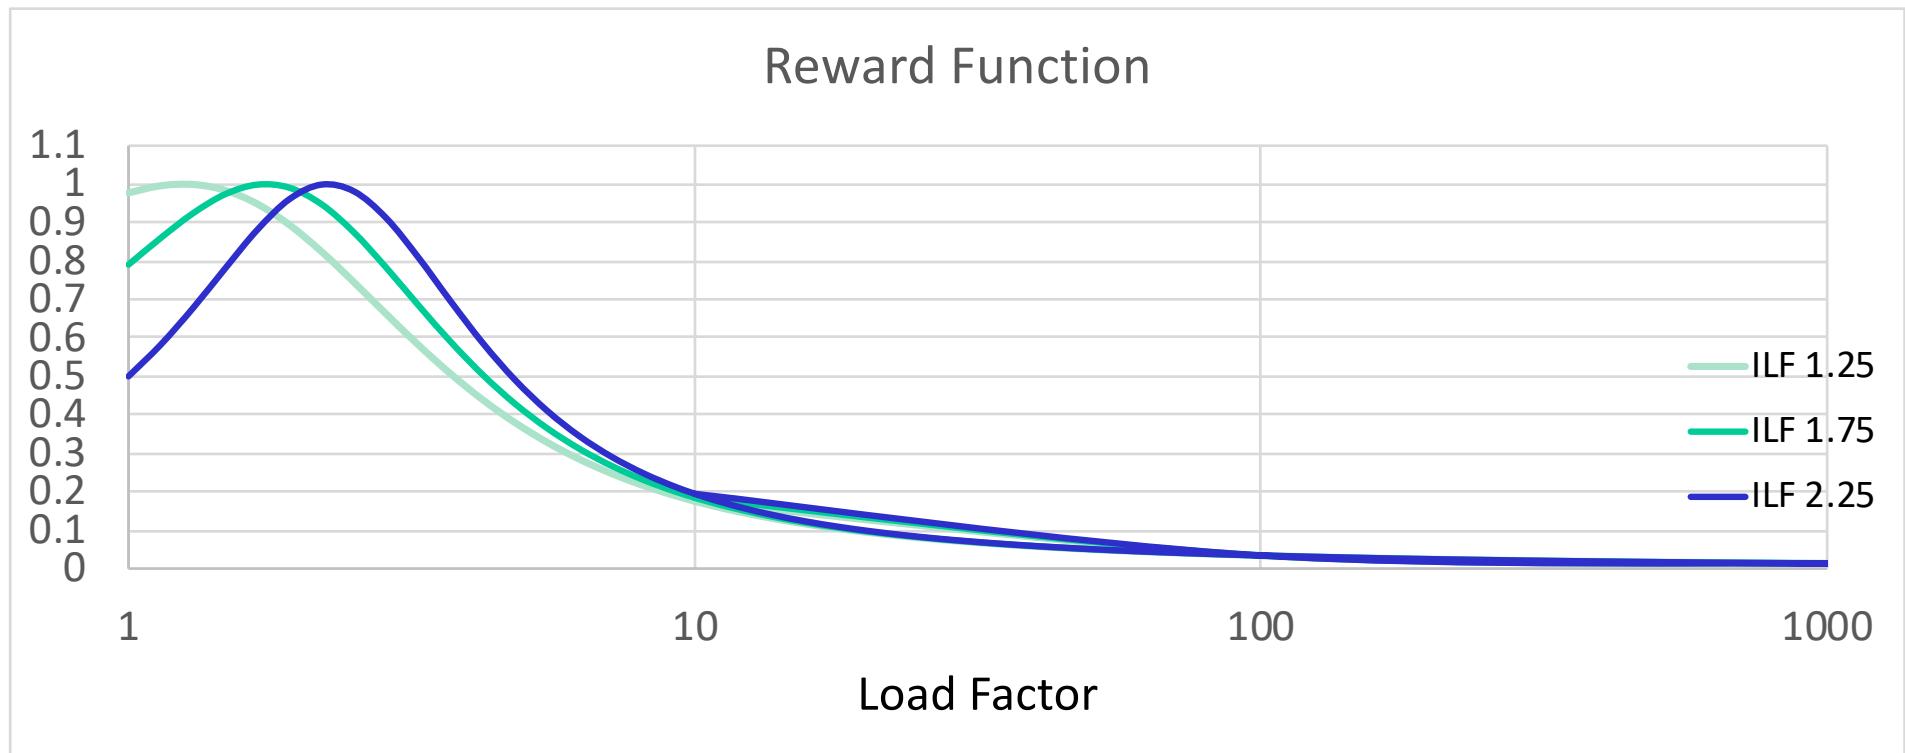
$$Reward(l) = \frac{1}{1 + (\log_2 [1 + \alpha(l - l^*)])^2}$$



- Maximized at ILF
  - Just above average population
  - Drops for smaller loads (too few) and larger loads (too crowded)

# Reward Function (cont.)

$$Reward(l) = \frac{1}{1 + (\log_2 [1 + \alpha(l - l^*)])^2}$$



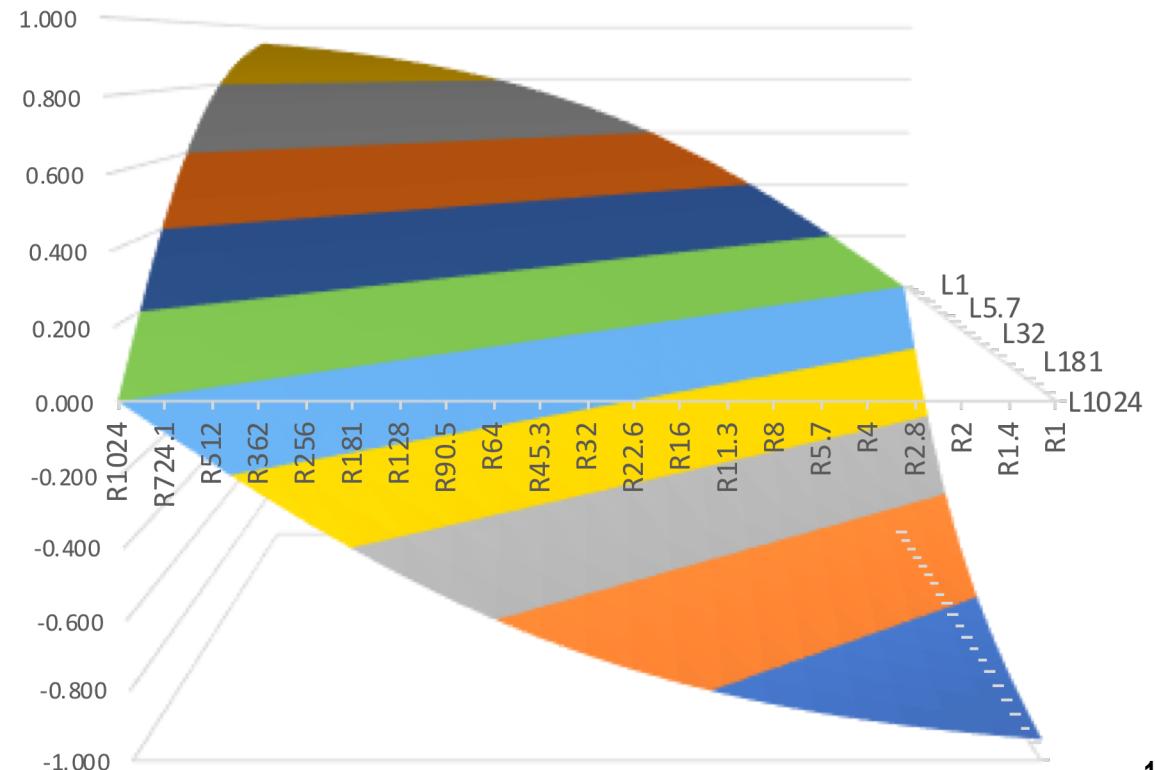
- Falls off gradually
  - $Reward(1000) = 0.0132$

# Computing Ideal Load Factor (ILF)

- Suppose node has count  $c_l$  and neighbor has count  $c_r$
- Compute *imbalance* as

$$\beta(c_l, c_r) = \frac{\sqrt{c_r} - \sqrt{c_l}}{\sqrt{c_l} + \sqrt{c_r}}$$

- Maximum +1  $c_r \gg c_l$
- Minimum -1  $c_l \gg c_r$



# Computing Ideal Load Factor (cont.)

- For node  $u$  with population  $p(u)$

$$\hat{\beta}(u) = \text{Avg}_{(u,v) \in E} [\beta(p(u), p(v))]$$

- Define ILF as

$$l^*(u) = 1.75 + 0.5 \cdot \hat{\beta}(u)$$

- Minimum 1.25

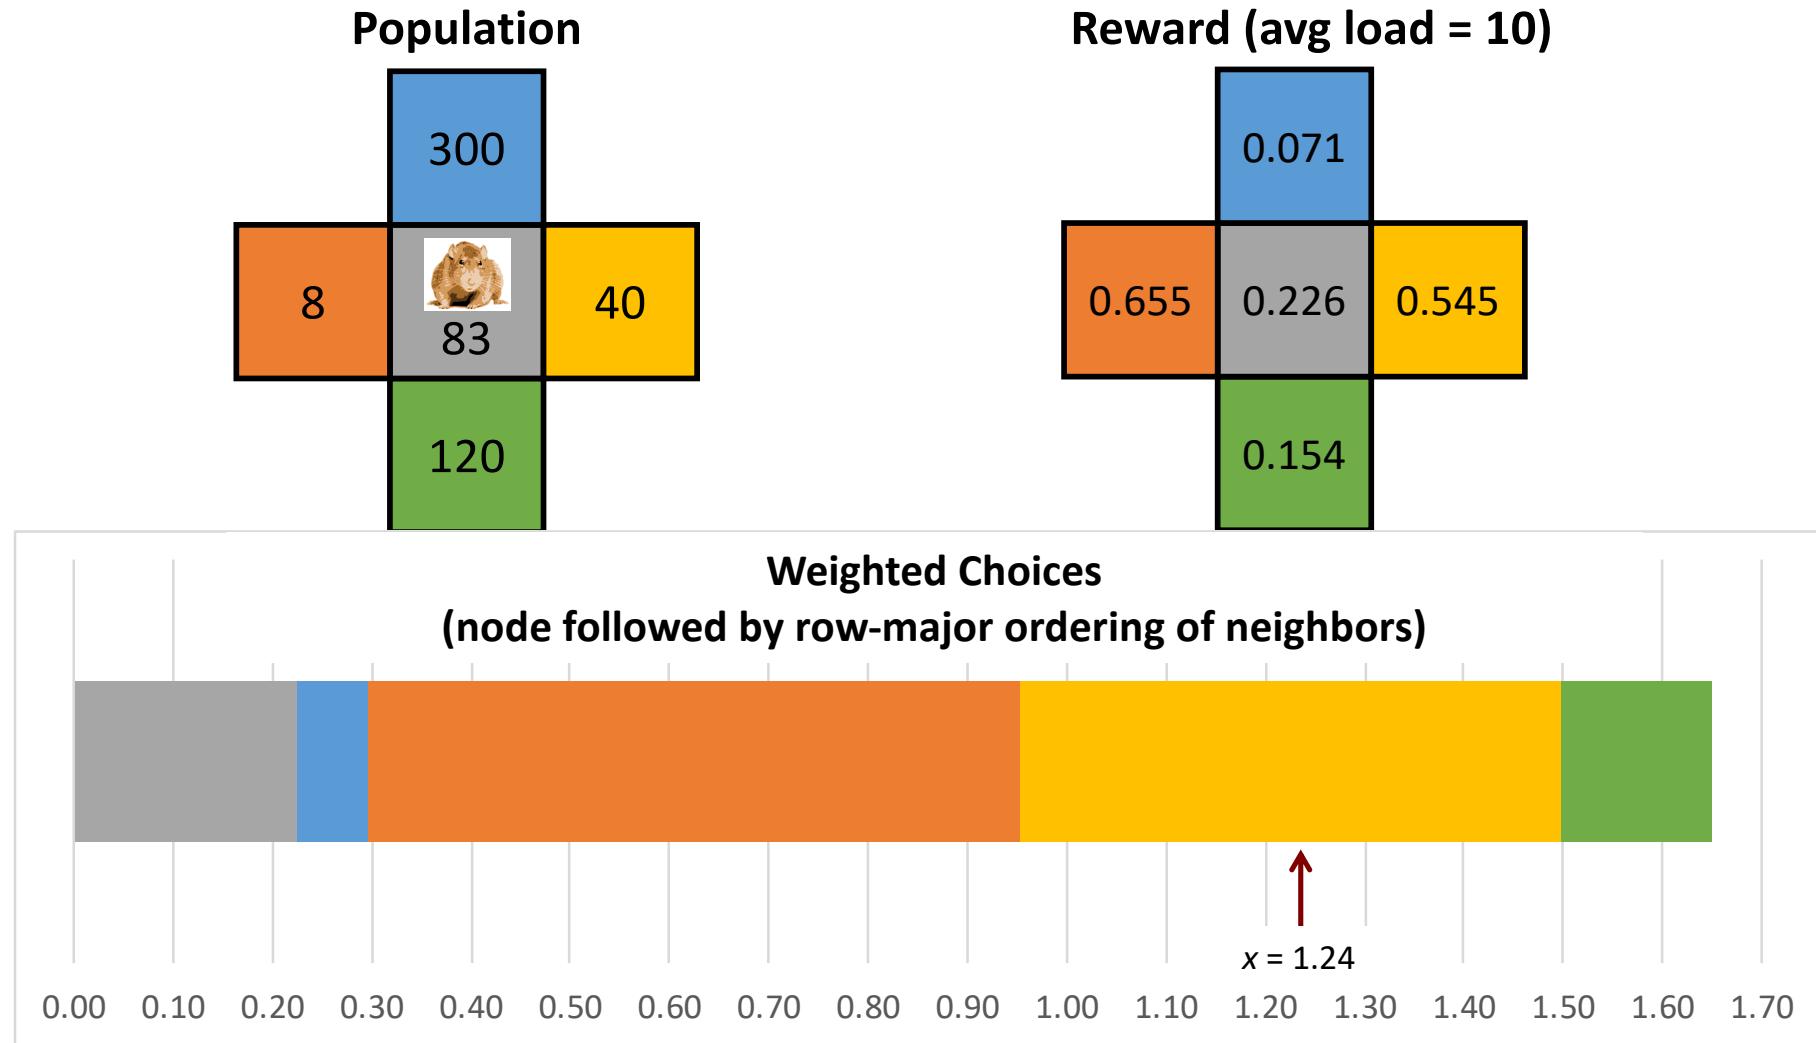
- When adjacent nodes much less crowded

- Maximum 2.25

- When adjacent nodes much more crowded

- Changes as rats move around

# Selecting Next Move



- Choose random number between 0 and sum of rewards
- Move according to interval hit

# Update Models

## ■ Synchronous

- Demo 3
- Compute next positions for all rats, and then move them
- Causes oscillations/instabilities

## ■ Rat-order

- Demo 4
- For each rat, compute its next position and then move it
- Smooth transitions, but costly

## ■ Batch

- Demo 5
- For each batch of  $B$  rats, compute next moves and then move them
- $B = 0.02 * R$
- Smooth enough, with better performance possibilities

# What We Provide

## ■ **Python version of simulator**

- Demos 1–2
- Very slow

## ■ **C version of simulator**

- Fast sequential implementation
- Demos 3–5: 36X32 grid, 11,520 rats
- Demos 6–11: 180X160 grid, 1,008,000 rats
  - That's what we'll be using for benchmarks

## ■ **Generate visualizations by piping C simulator output into Python simulator**

- Operating in visualization mode
- See Makefile for examples

# Correctness

## ■ Simulator is Deterministic

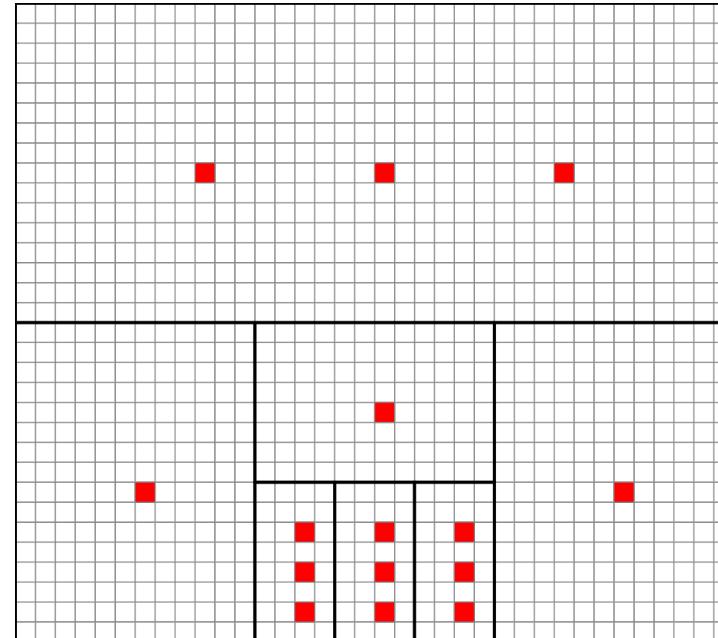
- Global random seed
- Random seeds for each rat
- Process rats in fixed order

## ■ You Must Preserve Exact Same Behavior

- Python simulator generates same result as C simulator
- Use `regress.py` to check
  - Only checks small cases
  - Useful sanity check
- Benchmark program compares your results to reference solution
  - Handles full-sized graphs

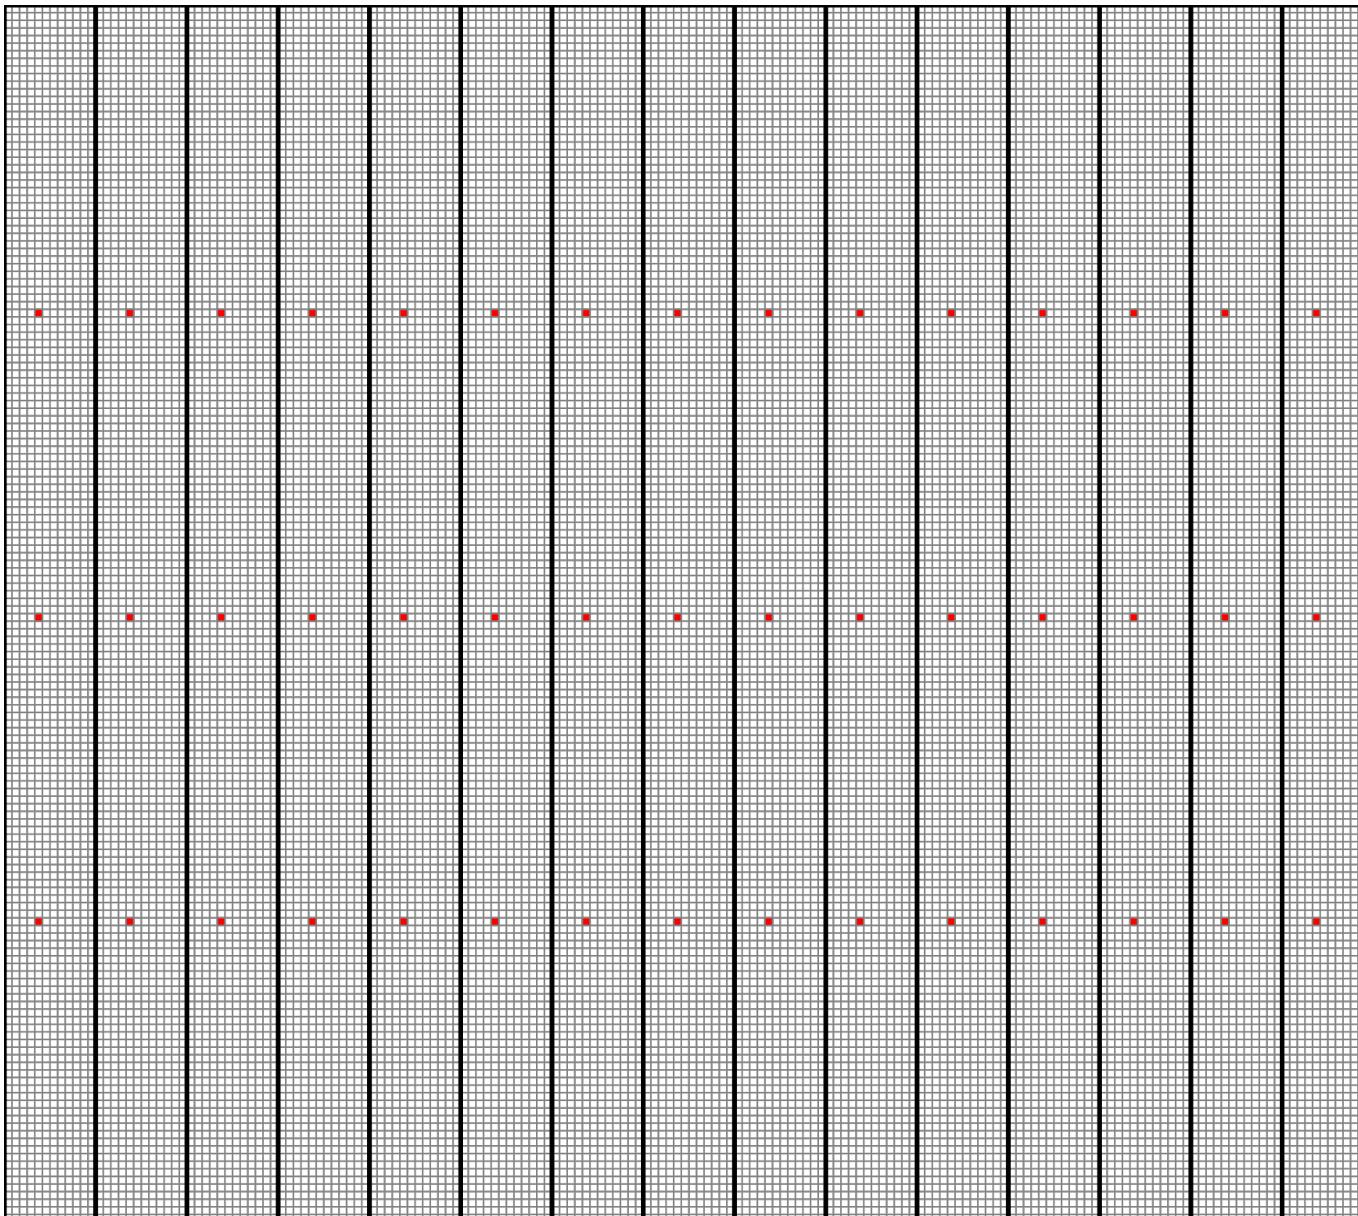
# Fractal Graph fracZ (Demos 1–5)

*Rats spread quickly within region  
More slowly across regions  
Hub nodes tend to have high counts*



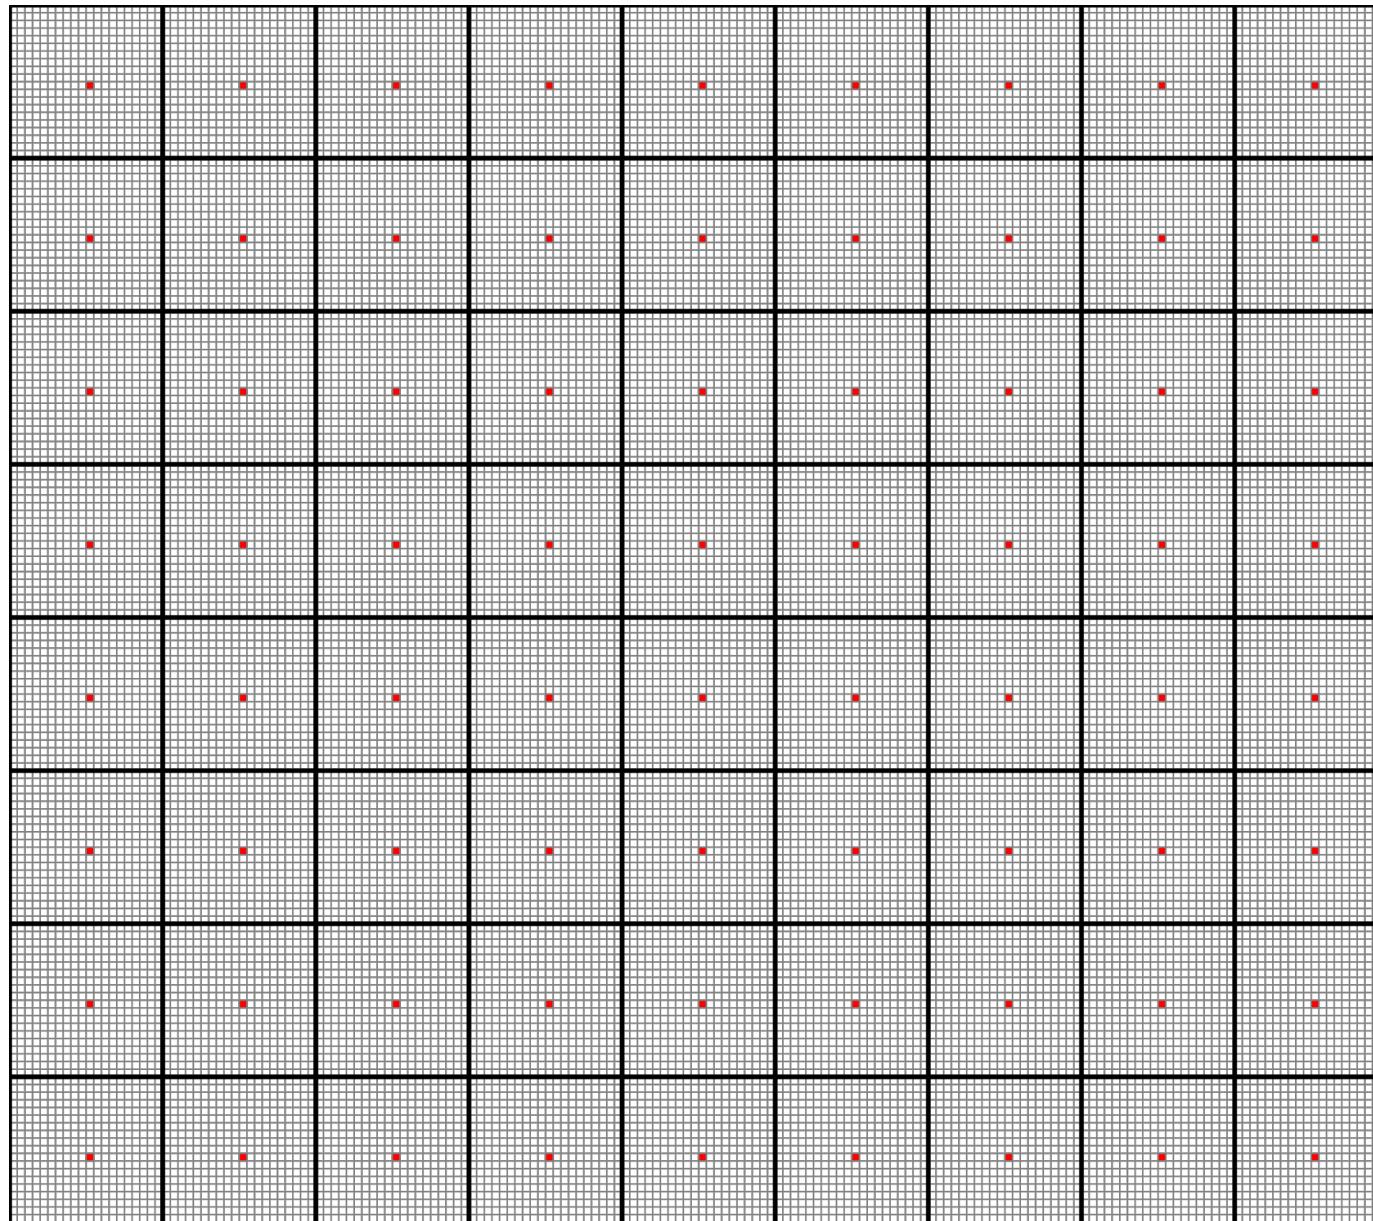
- **Base grid**
  - $W \times H$  nodes, each with nearest neighbor connectivity
- **Regions**
  - Recursively partition rectangles into two or three subrectangles
  - Leaves of tree form regions
- **Hubs**
  - Connect to every other node in region
  - Each region has one or three hubs

# Benchmark Graph UniA (Demo 7)



Nodes	28,800
Edges	286,780
Regions	15
Hubs	45
Max Degree	1,919

# Benchmark Graph UniB (Demo 8)



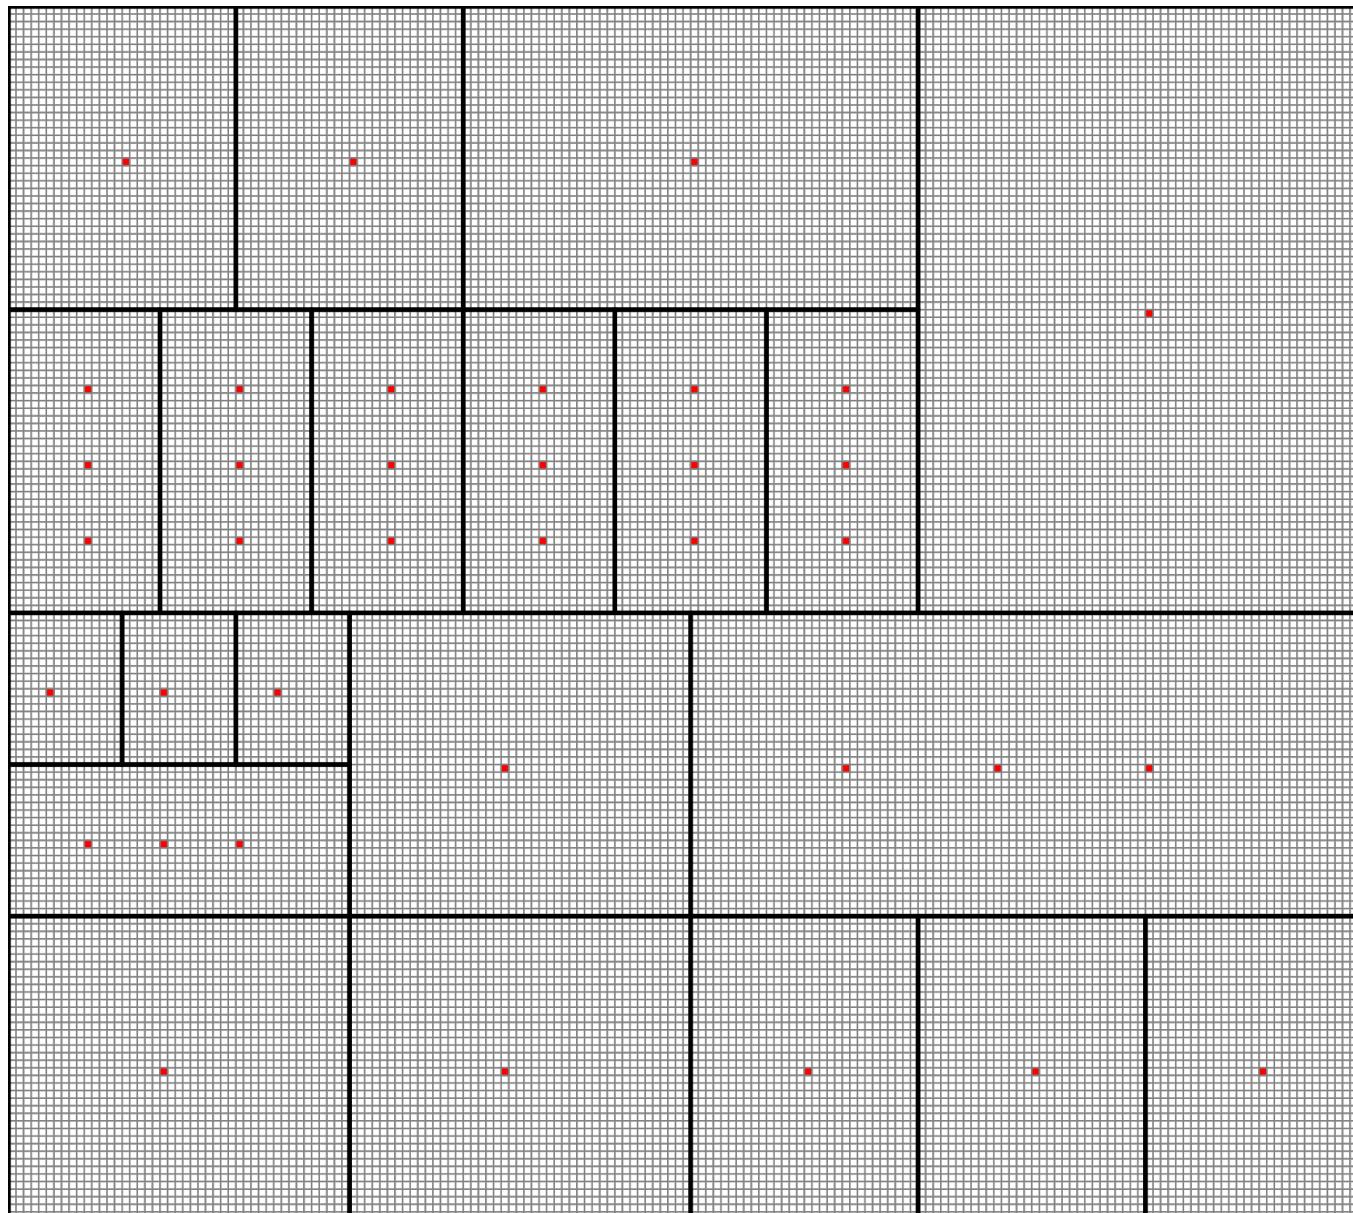
Nodes	28,800
Edges	171,400
Regions	72
Hubs	72
Max Degree	399

# Benchmark Graph FracC (Demos 9–10)



Nodes	28,800
Edges	187,612
Regions	30
Hubs	46
Max Degree	4,899

# Benchmark Graph FracD

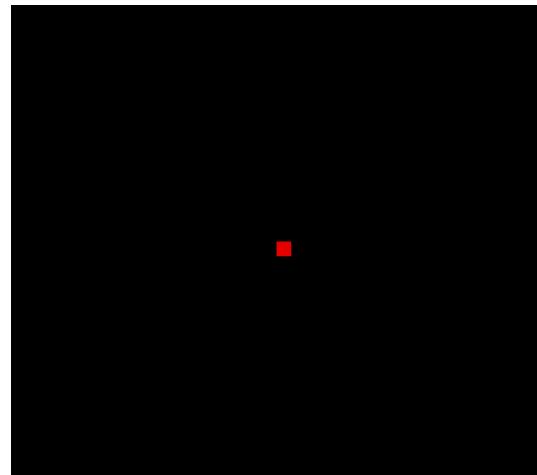


<b>Nodes</b>	<b>28,800</b>
<b>Edges</b>	<b>208,902</b>
<b>Regions</b>	<b>21</b>
<b>Hubs</b>	<b>37</b>
<b>Max Degree</b>	<b>4,899</b>

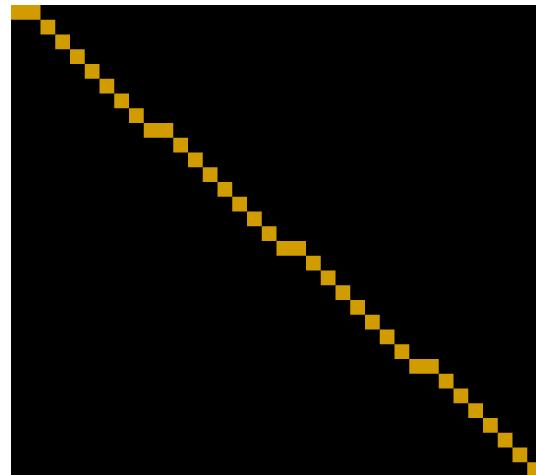
# Initial States (fracZ)

Center (r)  
~Demo 7

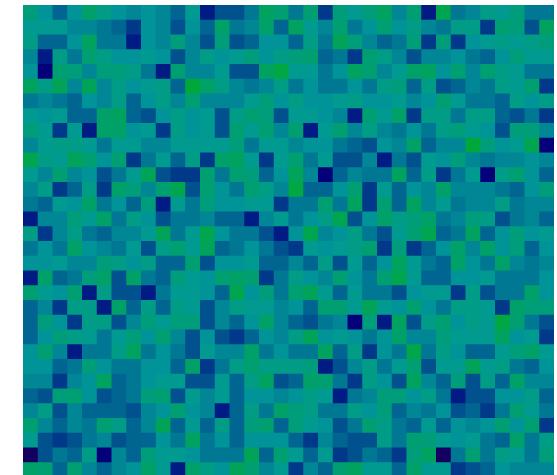
$t = 0$



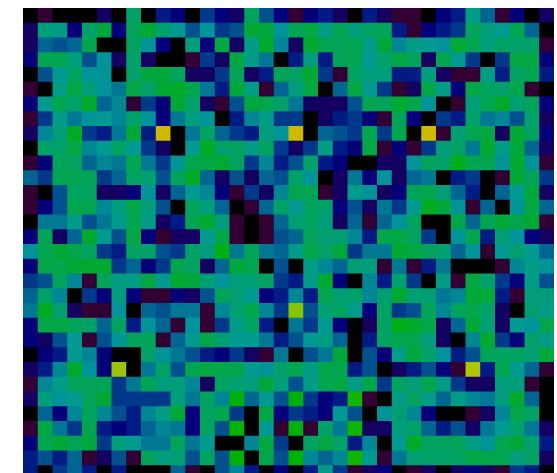
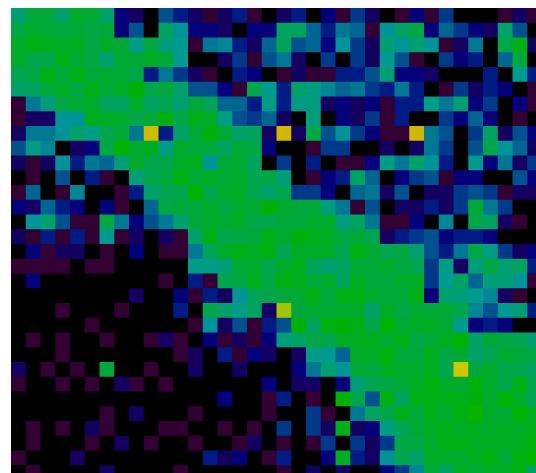
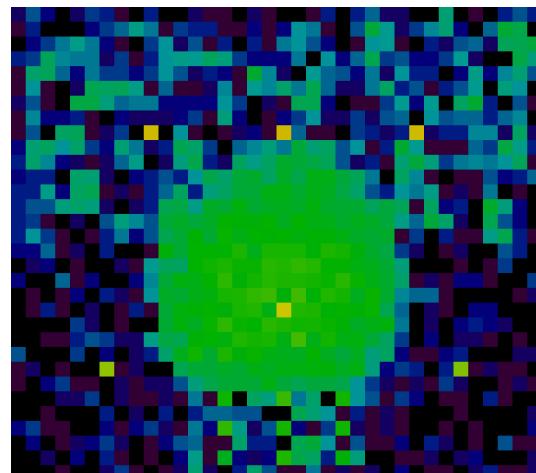
Diagonal (d)  
~Demo 9



Random (u)  
~Demo 10

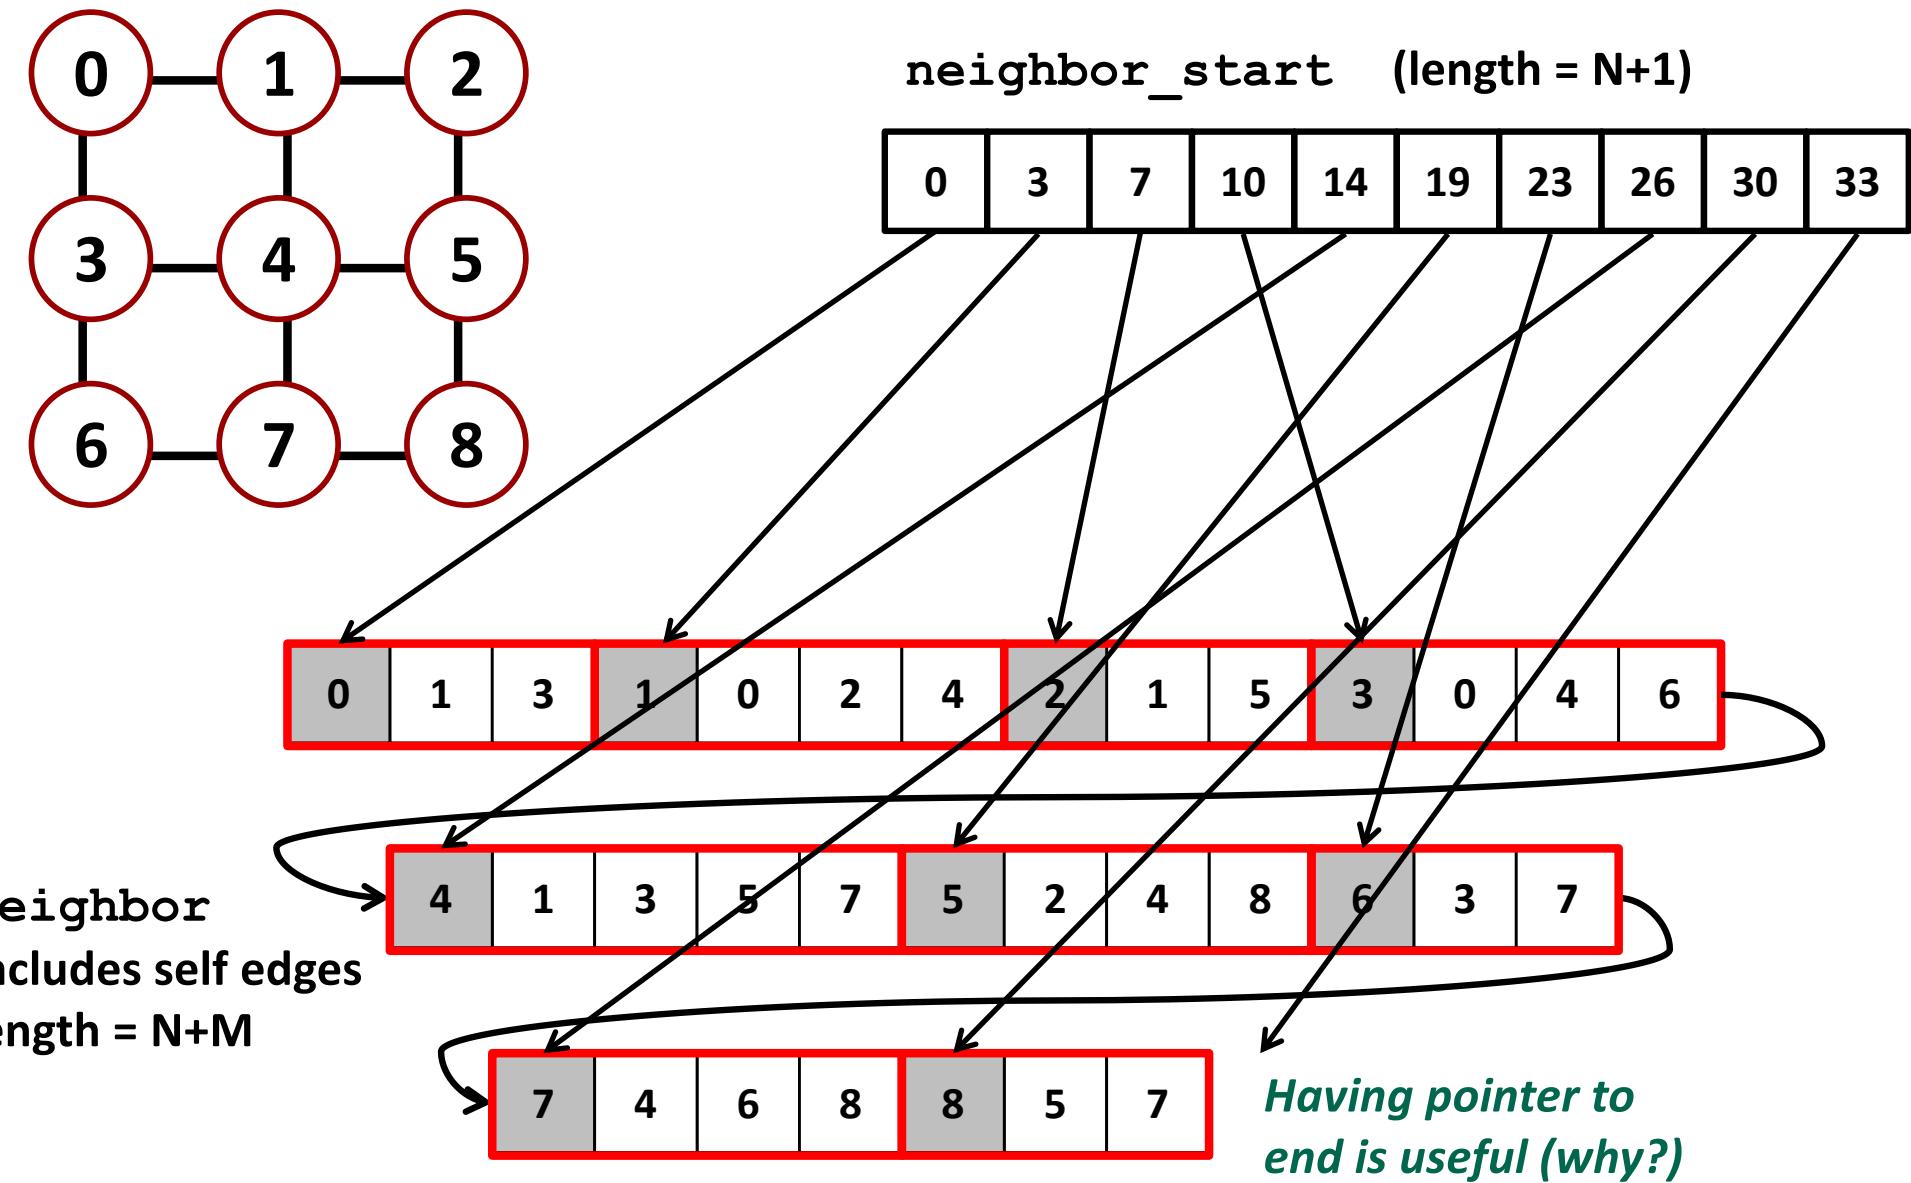


$t = 20$



# Graph Representation

N node, M edges



# Sample Code

- From sim.c
- Compute reward value for node

```
/* Compute weight for node nid */
static inline double compute_weight(state_t *s, int nid)
{
    int count = s->rat_count[nid];
    double ilf = neighbor_ilf(s, nid);
    return mweight((double) count/s->load_factor, ilf);
}
```

- Simulation state stored in `state_t` struct
- Reward function computed by `mweight`

# Sample Code

- From sim.c
- Compute reward value for all nodes

```
static inline void compute_all_weights(state_t *s) {  
    graph_t *g = s->g;  
    double *node_weight = s->node_weight;  
    int nid;  
    for (nid = 0; nid < g->nnode; nid++)  
        node_weight[nid] = compute_weight(s, nid);  
}
```

- Simulation state stored in **state\_t** struct

# Sample Code

- From sim.c
- Compute sum of reward values for node
- Store cumulative value for each edge
- Store total sum for later reuse

```
static inline void find_all_sums(state_t *s) {
    graph_t *g = s->g;
    int nid, eid;
    for (nid = 0; nid < g->nnode; nid++) {
        double sum = 0.0;
        for (eid = g->neighbor_start[nid];
             eid < g->neighbor_start[nid+1];
             eid++) {
            sum += s->node_weight[g->neighbor[eid]];
            s->neighbor_accum_weight[eid] = sum;
        }
        s->sum_weight[nid] = sum;
    }
}
```

# Sample Code

## ■ Compute next move for rat

```
static inline int fast_next_random_move(state_t *s, int r) {
    int nid = s->rat_position[r];
    graph_t *g = s->g;
    random_t *seedp = &s->rat_seed[r];
    double tsum = s->sum_weight[nid];
    double val = next_random_float(seedp, tsum);
    int estart = g->neighbor_start[nid];
    int elen = g->neighbor_start[nid+1] - estart;

    /* Find location by binary search */
    int offset = locate_value(val,
                               &s->neighbor_accum_weight[estart],
                               elen);

    return g->neighbor[estart + offset];
}
```

# Instrumented Code

- From sim.c
- Wrap major sections with instrumentation macros

```
static inline void find_all_sums(state_t *s) {
    graph_t *g = s->g;
    START_ACTIVITY(ACTIVITY_SUMS);
    int nid, eid;
    for (nid = 0; nid < g->nnode; nid++) {
        double sum = 0.0;
        for (eid = g->neighbor_start[nid];
             eid < g->neighbor_start[nid+1];
             eid++) {
            sum += s->node_weight[g->neighbor[eid]];
            s->neighbor_accum_weight[eid] = sum;
        }
        s->sum_weight[nid] = sum;
    }
    FINISH_ACTIVITY(ACTIVITY_SUMS);
}
```

# Running Instrumented Code

## ■ Demo 11

```
./crun-seq -g data/g-180x160-fracC.gph
  -r data/r-180x160-r35.rats -u b -n 50 -I -q
```

```
50 steps, 1008000 rats, 13.770 seconds
```

228 ms	1.6 %	startup
10677 ms	76.3 %	compute_weights
750 ms	5.4 %	compute_sums
2340 ms	16.7 %	find_moves
2 ms	0.0 %	unknown
13998 ms	100.0 %	elapsed

- Shows breakdown of where time spent
- See speedups of different parts of code
- Can instrument both your code & reference version

# Finding Parallelism

## ■ Sequential constraints

- Must complete time steps sequentially
- Must complete each batch before starting next
  - ILF values and weights then need to be recomputed

## ■ Sources of parallelism

- Over nodes
  - Computing ILFs and reward functions
  - Computing cumulative sums
- Over rats (within a batch)
  - Computing next moves
  - Updating node counts

# Performance Measurements

## ■ Nanoseconds per move (NPM)

- R rats running for S steps
- Requires time T seconds
- $NPM = 10^9 * T / (R * S)$
- Reference solution:
  - Average 290 NPM for 1 thread
  - Average 44.6 NPM for 12 threads
  - Speedups:
    - UniA 6.78
    - UniB 6.43
    - FracC 6.45
    - FracD 6.55
- Maybe you can do better!

# Performance Targets

## ■ Benchmarks

- 4 combinations of graph/initial state
- Each counts 16 points

## ■ Target performance

- $T$  = measured time
- $T_r$  = time for reference solution
- $T_r / T$  = How well you reach reference solution performance
  - Full credit when  $\geq 0.95$
  - Partial when  $\geq 0.60$

# Machines

## ■ **Latedays cluster**

- 16 worker nodes + 1 head node
- Each is 12-core Xeon processor (dual socket with 6 cores each)
- You submit jobs to batch queue
- Assigned single processor for entire run
- Python script provided

## ■ **Code Development**

- OK to do code development and testing on other machines
- But, they have different performance characteristics
- Max threads on GHC cluster = 8
- Code should run for any number of threads (up to machine limit)

# Some Logos



**GraphChi**: Going small with GraphLab

